KUKU AI Problem 1: Find the Promos

Measured solution summary

Approach

A reusable Python pipeline parses all timestamped screenplay events, generates bounded event windows , scores hook, cliffhanger, intensity, standalone clarity, and length, then applies hard validation before ranking.

Measured run

50 episodes | 7026 candidates | 1781 valid | 63.25 seconds

The source pack remains private and is excluded from version control.

Validation and ranking

Independent duration, bounds, ordering, gap, boundary, hook, cliffhanger, and diversity checks prevent AI scores from authorizing invalid promos.

Results and limitations

The run produced 86 ranked selections. No gold or human review data was supplied, so precision and recall are unavailable. The weakest selections can still be coherent but overly resolved; human review is recommended before publication.

AI disclosure

Tools used: GitHub Copilot and optional Google Gemini API. Example mistake: the initial heuristic treated a resolved statement containing 'leave' as a cliffhanger; closure detection was added after inspecting the actual top picks.